

Executive Report: NovaTech Revenue Intelligence Dashboard

Connecting marketing, sales and customer support in one view for the revenue team

Date	13 September 2026	Prepared for	Sarah Chen, VP of Revenue
Written by	Eric Jang, Business Intelligence Analyst	Purpose	Revenue Intelligence Dashboard review and recommended actions

1. Technical terms on first use

The following terms are used throughout this report and are defined here on first use.

Term	Meaning in this report
SPICE	SPICE (Super-fast, Parallel, In-memory Calculation Engine) is Amazon Quick’s in-memory data store. It holds a copy of each dataset so dashboards load quickly without querying the source systems, and it is refreshed weekly.
ROI (return on investment)	ROI measures how much a campaign earned back relative to its cost. It is attributed revenue minus campaign spend, expressed as a percentage of spend, so a negative ROI means the campaign earned back less than it cost.
Attributed Revenue	Attributed revenue is the revenue that the marketing platform credits to a specific campaign interaction. It is a marketing measure and differs from the deal revenue booked in the CRM (customer relationship management system).
SLA (service level agreement)	An SLA is a committed target for how quickly Support must resolve a ticket, normally set by priority level, such as a maximum resolution time for critical tickets.
Account ID	Account ID is the unique customer identifier (for example, ACCT-041) used in the CRM, marketing and support systems. It is the shared key that connects the three sources for the same customer.
Quick Chat	Quick Chat is the natural-language assistant in Amazon Quick that answers plain-English questions about dashboards and data. Its answers are referred to as “Q” in this report.
Topic (AWS Quick)	A Topic is a business-language layer on top of the datasets. It gives fields friendly names, synonyms and definitions, hides ambiguous fields and defines how datasets relate, so that Quick Chat interprets questions the way the business does.

2. Executive Overview

NovaTech’s revenue team currently relies on three separate systems for sales, marketing and customer support. This fragmentation requires the team to combine reports manually each Monday, consuming roughly an hour every week while still leaving cross-functional questions difficult to answer. The Revenue Intelligence Dashboard addresses this problem by bringing the customer journey into three connected views, **Marketing Funnel**, **Sales Pipeline** and **Customer Health**, with natural-language access through Quick Chat.

The dashboard has been designed around the questions that matter most to the revenue team: *which marketing investments convert, where sales performance is strongest or slowing, and which customers or product areas represent the greatest retention risk*. It deliberately focuses on current and historical performance rather than forecasting, consistent with the requirement for weekly rather than real-time data. The intended result is a single decision-making environment that links acquisition, revenue and customer experience, and its greatest value is enabling NovaTech to identify where revenue is being lost, delayed or put at risk and to assign a clear business response.

2,240	–90.88%	\$707.20K	63.13%	3,000	22.8%
marketing leads	overall campaign ROI	won revenue	deal win rate	support tickets	negative sentiment

3. Data Strategy

The solution combines three source datasets: **499 CRM deals**, **2,240 marketing leads** and **3,000 support tickets**. All three share Account ID, which provides the common identifier needed to connect marketing activity, sales outcomes and support experience. NovaTech’s CRM contains 85 accounts, while the marketing and support systems contain an additional 15 account IDs (ACCT-101 to ACCT-115, covering 150 leads and 204 tickets) that do not appear in the CRM. These records remain useful within the marketing and support views but cannot be linked to recorded deal revenue.

Before the dashboard was built, each dataset was loaded into SPICE and checked against the data dictionary, the reference document that describes every field. The checks confirmed record counts, date ranges and missing values, and dates were set up to be recognised as dates rather than plain text. Three business measures were then added: **days to close** (the time from a deal being opened to being closed), **campaign ROI** and **ticket resolution time** (the hours from a ticket being raised to being resolved). The checks also highlighted data-quality considerations that should continue to be monitored: 59 tickets have

no resolution date because they are still open, 59 tickets carry a blank sentiment value that the data dictionary does not document, and a small number of deal and ticket IDs are repeated across different accounts.

To answer questions that cross all three systems, a fourth, **combined dataset** places each customer's marketing, sales and support information side by side, matched by Account ID. Support tickets form the base of this dataset, known as its "anchor", and each ticket carries a one-line summary of that customer's deals and marketing activity. Every ticket is kept even without a matching summary, an approach known technically as a "left join".

The summaries were created deliberately before the data was combined. Because one customer can have many deals, leads and tickets, lining up every record against every other record for the same customer would have turned 499 deals into about **63,420 overlapping rows** and made revenue appear roughly 143 times larger than it really is. Summarising first means that each ticket is counted once and each customer's revenue is shown once rather than added up ticket by ticket. Headline figures such as won revenue, win rate and campaign ROI are still calculated from their original source, so the combined dataset is used only where the revenue team needs to compare marketing, sales and support for the same customers.

4. Dashboard Design Rationale

Each view maps directly to a section of the requirements brief. Chart types follow a simple rule: KPI (key performance indicator) cards for headline measures, a funnel for stage progression, bar charts for comparisons across campaigns, channels, products and regions, line charts for trends over time, and tables where account-level detail matters.

View	Questions it answers	Key visuals	Filters and interactions
Marketing Funnel	It shows how campaign investment turns into leads, responses and closed-won outcomes, which channels convert and which campaigns pay for themselves.	KPI cards show leads, response rate, conversion, spend, attributed revenue and ROI. A funnel shows each stage, and charts compare channels, campaigns, and spend against revenue.	Users can filter by campaign, channel, date range and customer segment.
Sales Pipeline	It shows won versus lost deals and why, revenue by segment and product, win rates by region, rep and product, days to close and its trend, and average deal value.	KPI cards show won revenue, win rate, average won deal and days to close. Charts compare outcomes, loss reasons, revenue by product and size, and win rate by region and rep, with days to close by quarter.	Users can filter by region, manager, product line, date and deal stage, and clicking a product or region bar filters the view.
Customer Health	It shows which accounts file the most tickets and whether they are the largest revenue accounts, resolution time by priority, product-area load, sentiment and accounts at risk.	KPI cards show tickets, resolution time, negative sentiment, and open and critical tickets. Charts show resolution by priority, tickets by product area and sentiment, and an account risk view uses the combined dataset.	Users can filter by account, priority, product area, support tier and region, and jump from an account to its deals.

Two further design choices are deliberate. First, "segment" means three different things in the source systems (the marketing lead segment, the CRM company-size tier and the support service tier), so each is labelled separately to avoid misleading comparisons, and Account ID is used as the reliable link between views. Second, forecasting and print-ready layouts were left out, in line with the brief: the dashboard is an interactive weekly tool for the revenue team rather than a board pack.

5. How Topic configuration affected Q accuracy

The Topic was designed to make natural-language questions more reliable by giving business-friendly names to fields, adding synonyms, clarifying concepts such as *deal revenue* and *conversion rate*, and excluding ambiguous fields that could cause Quick Chat to select the wrong definition. For example, a customer's own annual company revenue is deliberately distinguished from NovaTech's deal revenue, while resolution time, campaign response and conversion have explicit business meanings.

To measure the effect, the same baseline questions on won revenue, the highest-converting channel and resolution time by priority were asked before the Topic was configured and re-asked word for word afterwards. After configuration, Q's figures were consistent with the dashboard: **\$707,201** in won revenue, a **48.99%** conversion rate for Direct Mail, and **56.13 hours versus 58.92 hours** for critical and low-priority tickets. For resolution time, the figures did not change between the two runs, but Q's narrative interpretation of the gap did, showing that the Topic shapes how Q explains a result as well as which field it selects. It did not remove the limits described in Section 7. Quick Chat should therefore be treated as a fast exploration layer rather than the final system of record, meaning the official source of reported figures, and its answers should be verified against the relevant dashboard view.

6. Key Insights with Recommended Actions

Insights	Implications	Recommended Actions	Owner + Estimated timeframe
1. Marketing investment is not paying back. \$12.36M of campaign spend produced \$1.13M of attributed revenue (ROI –90.88%), and 412 leads closed-won, an 18.39% lead-to-closed-won conversion rate. NovaEdge Awareness consumed about \$1.52M for about \$34K of attributed revenue, with a 6.2% response rate and about 5.1% conversion. NovaPulse Launch is the strongest campaign by closed-won volume, with about \$394K of attributed revenue.	Budget is concentrated in a programme that does not convert, leaving less for campaigns that do.	Reduce or pause NovaEdge Awareness and model a controlled reallocation towards stronger campaign and channel combinations. Confirm with Finance first that attributed revenue captures the right economic value.	Marketing, with VP of Revenue approval Timeframe: The pause is decided within two weeks, the model is ready within 30 days, and approval comes before the next campaign cycle.
2. Channel efficiency varies substantially. Direct Mail converts 48.99% of its leads, considerably ahead of Partner Referral at 31.1%.	Channel mix, not only total spend, decides how many leads become customers.	Use channel conversion as a core input to the reallocation model and re-measure the results after one campaign cycle.	Marketing Timeframe: Inputs are due within 30 days, with results re-measured after one quarter.
3. Deals are closing more slowly. The pipeline holds 499 deals with a 63.13% win rate, \$707.20K in won revenue and an average won deal of about \$2,245. Enterprise and Large customers contribute most revenue (about \$319K and \$272K), and NovaPulse Professional leads products (about \$209K). Average days to close reached 90.1 in Q1 2025 against 51.5 in Q1 2024, an increase of about 75%. No Decision Made and Poor Product Fit are the leading loss reasons (43 deals each).	Longer cycles delay revenue and raise the risk of stalled deals. Q1 2025 covers January only, so the trend needs confirming.	Introduce a 60-day opportunity review with alerts for deals that stay too long in one stage, strengthen early product-fit qualification and monitor deal-cycle length weekly.	Sales managers Timeframe: Weekly monitoring starts at the next Monday review, alerts go live within 30 days and the qualification checklist within 60 days.
4. Urgent tickets are not resolved faster. Across 3,000 tickets there are 50 critical tickets, 59 open tickets and 22.8% negative sentiment. Average resolution is about 58 hours and changes little by priority: 56.13 hours for critical versus 58.92 hours for low-priority tickets.	Customers with business-stopping issues wait almost as long as those with minor ones.	Set SLAs by severity, with a tighter critical-ticket target, and track attainment on the Customer Health view.	Customer Success and Support Timeframe: Draft SLAs are due within 30 days and should be in effect within 60 days.
5. Two product areas drive most severe issues. Authentication and Data Pipeline account for 279 of the 450 high-severity tickets (high and critical priority), or 62%, despite representing only 29.7% of total ticket volume.	These areas can block access or disrupt core workflows, raising negative sentiment and churn risk.	Treat both areas as Priority 0 (P0, highest-priority) reliability initiatives with shared ownership and clear targets.	Engineering, with Support Timeframe: Owners are named within two weeks and an improvement plan agreed within 60 days, with quarterly reviews.
6. The largest accounts are also the most demanding. YieldMax Software, the largest account by won revenue (\$40,722), also files the most tickets (334).	Revenue and support load overlap in the accounts NovaTech can least afford to lose.	Review high-risk accounts regularly with Engineering, Support and Revenue leadership.	Customer Success Timeframe: The first review takes place within two weeks, followed by monthly reviews.

Operations should refresh the dashboard weekly and make it the primary Monday revenue review, using Quick Chat for exploratory questions rather than replacing governed dashboard metrics. Timeframes are estimates to confirm with owners.

7. Where analysis agreed or disagreed with the dashboard

Where did Q agree with the dashboard? Q matched the supporting visuals for **\$707,201** in won revenue, Direct Mail's **48.99%** conversion rate, and average resolution times of **56.13 hours** for critical tickets and **58.92 hours** for low-priority tickets. It also found that every campaign's spend exceeded its attributed revenue and reproduced the top-10 ticket-volume ranking.

Where did Q disagree or struggle, and why? Q struggled when the requested calculation was not available in a visual or when the Topic lacked the required relationship. It substituted total deal value for average deal size, could not combine the 30-day account ticket threshold with per-deal values, and could not connect Partner Referral sourcing to deal outcomes. Its narrative interpretation of the resolution-time gap also changed even though the figures did not.

When would you use Q versus the dashboard? I would use Q for quick questions that map directly to defined fields, calculations and relationships. I would use the dashboard to monitor performance, apply filters, inspect trends and verify Q's response against a visible measure. For questions outside the current semantic model (the Topic's set of definitions), I would first add or correct the required relationship, calculation or visual, then retest Q before using the result for a decision.